

Lifelong Language Learning with Adapter based Transformers

Anonymous EMNLP submission

Abstract

Continuous learning is vital in real-world applications of natural language processing, where interaction with ongoing streams of tasks and language is necessary for real-world NLP applications. When forced to adapt to new tasks and inputs, language models experience catastrophic forgetting. The current generative replay-based algorithms are not scalable to many tasks, and their performance may degrade from a change in the task order. In this paper, we propose a model based on network growth - a pre-trained Transformer with Adapter modules for each task - that sequentially learns new NLP tasks in various domains and prevents catastrophic forgetting without retraining the model from scratch. We train and maintain light weight adapter modules sequentially for each task. Without increasing network growth by more than 15% and avoiding replay and task order biasness, the current design allows us to increase average task accuracy by 1.3% over the baseline models.

1 Introduction

Humans can continuously accumulate, develop, and transfer knowledge and skills throughout their lifetimes, giving rise to lifelong learning principles. Continuous learning is crucial in real-world natural language processing applications, where computer systems must interact with ongoing streams of data and language across time. Isolated learning is currently the most used paradigm in machine learning. In isolated learning, the model experiences catastrophic forgetting or interference due to non-stationary data distribution that biases the model, making it unable to remember the information it has previously learned on a stream of tasks joined to be trained sequentially.

In comparison, continual learning focuses on the ability of the model to learn continuously and adaptively over time, which allows it to learn new information without forgetting the past knowledge.

In this work, we focus on lifelong language learning (LLL) on a continuous stream of NLP tasks. The performance of LLL is typically viewed as having an upper bound provided by multi-task learning. There is still a performance gap between other frameworks for LLL (Sun et al., 2019; Huang et al., 2021; Kanwathara et al., 2021) and multi-task learning. In these previous works, continual learning is maintained through generative replay which limits its application to large number of tasks and the model performance changes as the task order changes (Sun et al., 2019).

We enhance the current LLL strategies by proposing a novel approach - Adapter based Transformer - a dynamic architecture based on network growth. Particularly, we propose to use task-specific adapters for each task, which equips the framework to learn new tasks while retaining information about older tasks and thus, avoid catastrophic forgetting.

Our main contributions are :

- We present Adapter based Transformers, our framework is space efficient. Due to the adapters being lightweight, very little extra memory is utilized.
- Number of tasks to be trained need not be known in advance, we can always learn new tasks by adding new adapter modules.
- The performance of our framework is independent of the order of the tasks.

2 Background and Related Work

2.1 Continual Learning

Among existing continual learning approaches for LLL, replay-based methods (Sun et al., 2019) and regularization-based methods (Huang et al., 2021) have been widely applied to NLP tasks to enable large pre-trained models to acquire knowledge from streams of textual data without forget-

ting the already learned knowledge. LAMOL (Sun et al., 2019) is a data-based LLL approach that simultaneously learns to solve a new task, while generating pseudo samples for previous tasks to train alongside the new task. A single model is used here, and no extra generator is used. Rational LAMOL (Kanwatchara et al., 2021) is an enhancement on LAMOL. This framework applies freezing to critical components, that are identified by rationales, which are part of input texts that best explain the prediction or class labels, in transformer based language models, to maintain previously learned knowledge while being trained on a new task. This is done as pseudo-sample generation may not be sufficient to prevent catastrophic forgetting. In Information Disentanglement based Regularization (Huang et al., 2021), the framework focuses on how to generalize models to new tasks, rather than just focusing on preserving knowledge from previous tasks. It uses a multi layer encoder for a given sentence, that outputs hidden representations which contain generic as well as task specific information, and two disentanglement networks to extract the generic and specific representations. While training on new tasks, the model regularizes both these representations to different extents, to better remember previous knowledge as well as transfer to new tasks.

2.2 Adapters

In a large pre-trained model with parameters Θ , adapters (Houlsby et al., 2019) are neural modules with a limited number of newly added parameters Φ . While keeping Θ constant, the parameters Φ are learned on a target task; as a result, Φ learn to encode the task-specific representations in the pre-trained model’s intermediate layers. When compared to the amount of parameters in a pre-trained model, the number of parameters in adapters is parameter-efficient, comprising only 1% to 3% of those parameters. Due to their modularity and small size, (Pfeiffer et al., 2020), they accelerate training iterations and are shareable and composable. Different adapter architectures are tested in the (Houlsby et al., 2019) experiment, and the results empirically demonstrate the effectiveness of a two-layer feed-forward neural network with a bottleneck.

A schematic representation of the adapter module is provided in 1.

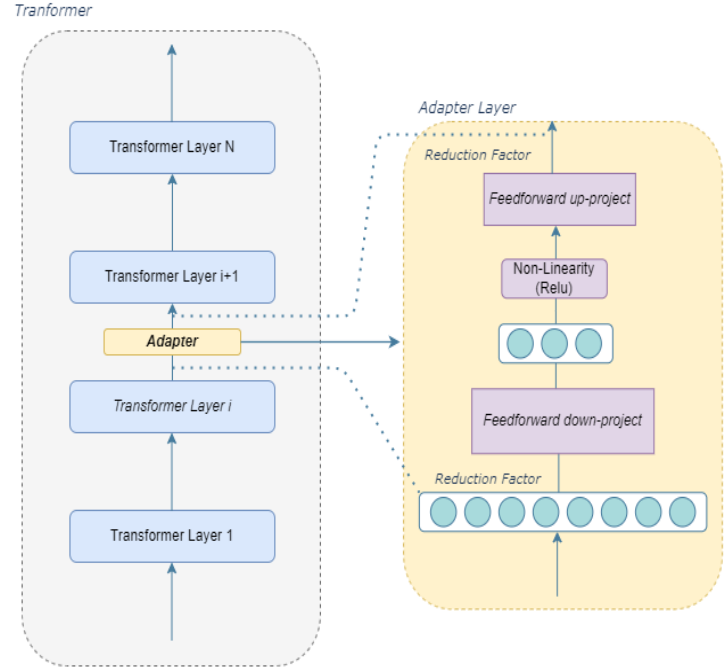


Figure 1: Adapter Model

3 Adapter based Transformers

We formalize our problem as a lifelong learning problem on a set of NLP tasks $\{T_1, \dots, T_N\}$ taken sequentially, where the number of tasks may be unknown. We propose Adapter based Transformers, a framework that contains adapter modules specific to each task. Our framework sequentially learns new NLP tasks in different domains and avoids catastrophic forgetting by storing relevant information of older tasks in respective task-specific adapters. We used a pre-trained GPT-2 model (Radford et al., 2018) for the transformer. Whenever a new task comes, we add two new adapters specific to the task, to each layer of the transformer. The task-specific adapter is then trained, keeping the weights of the underlying language model and non-task specific adapters frozen. Therefore, the model remembers previous tasks perfectly, while simultaneously being able to learn the new task. Since adapters have smaller number of parameters as compared to the original network, the model size grows slowly as more tasks are added.

For each new task T_i , the parameters of the pre-trained transformer as well as older task-specific adapters are kept frozen, and the adapter corresponding to the new task is trained. The parameters Φ_i of Adapter A_i are randomly initialized, and A_i is then trained on the new task T_i . This enables our

framework to retain knowledge of older tasks while simultaneously learning new tasks. All methods use the smallest pre-trained GPT-2 model (Radford et al., 2018) as the LM.

4 Experimental Setup

4.1 Data Formatting

Inspired by the protocol used in decaNLP (McCann et al., 2018) and LAMOL (Sun et al., 2019), samples from the datasets we use are framed into a SQuAD-like scheme, which consists of context, question, and answer. Special tokens are also added: ANS is inserted between question and answer. As the context and question are known during inference, decoding starts after inputting ANS.EOS is the last token of every example. Decoding stops when EOS is encountered.

4.2 Dataset Details

We conducted experiments on the following datasets to evaluate our proposed framework:

- Stanford Sentiment Treebank(SST) (Socher et al., 2013) - which is a sentiment analysis task consisting of sentiments (binary - positive or negative) corresponding to a movie review.
- Semantic Role Labeling (SRL) (He et al., 2015) - which is a role labeling task where Wikipedia domain of QA-SRL 1.0 is used, and the task is to assign labels that assigns semantic meaning to a phrase or sentence.
- Goal-oriented dialogue (WOZ) (Mrkšić et al., 2017) - which is a reservation task from English Wizard of Oz restaurant, and it comes with predefined information that will assist an agent to make a reservation for the customer.

Table 1 contains a summary of the datasets, dataset sizes, and metrics.

Dataset	Train	Test	Metric
SST	6920	1821	EM
SRL	6414	2201	nF1
WOZ	2536	1646	dsEM

Table 1: Details of dataset size and metrics. Here, nF1 indicates normalized F1 score; EM indicates exact match between texts; dsEM indicates turn-based dialogue state exact match

4.3 Baseline

LAMOL: We wanted to compare the best version of LAMOL (Sun et al., 2019) with our model; hence LAMOL’s best-performing parameters were used. We have set $k = 20$ in top-k sampling and $\lambda = 0.2$ for the weight of the LM loss. The learning rate is set up as $6.25e - 05$ and is scheduled to be linear with a warmup. Each task is trained for five epochs with loss as a summation of Question Answering (QA) and Language model (LM) loss and the optimizer is AdamW (Loshchilov and Hutter, 2019).

Adapter: To compare it more accurately, we have used almost similar parameters. We have set $k = 20$ in top-k sampling and $\lambda = 0$ for the weight of the LM loss. The learning rate is set up as $6.25e-05$ for all the tasks and is scheduled to be linear with a warmup. Each task is trained for 12 epochs with loss as only Question Answering (QA) loss and the optimizer is AdamW (Loshchilov and Hutter, 2019).

5 Results

This section reports the performance of Adapter-Transformers and compares it with LAMOL (Sun et al., 2019) as the baseline.




Figure 2: Performance Comparison between proposed Adapter-Transformers and baseline for three tasks

The proposed adapter model sees a 1.3% increase in average task accuracy as compared to baseline, seen in Table 3 and Figure 2 while avoiding replay and task order biasness.

As our proposed model is light-weight and avoids replay, we observe in Table 4 that it has significantly faster training time compared to baseline, with less than a 15% increase in network growth, as can be seen in Table 2. Figure 3 shows the proposed model’s loss, the loss decreases significantly for each of the three tasks.

	#Adapter-Transformer	#LAMOL	Increase	Increase %
1 layer	8,072,320	7,087,104	985,216	13.9
12 layers	96,867,840	85,045,248	11,822,592	13.9
Total(Layerwise+Fix)	136,254,720	124,432,128	11,822,592	9.5

Table 2: Summary of Parameter increase in proposed model vs baseline

Dataset	Metrics	Adapter	LAMOL
SST	EM	90.88	90.94
SRL	nF1	67.96	68.38
WOZ	dsEM	89.34	85.75

Table 3: Details of averaged metric scores for our proposed Adapter-Transformer method and baseline LAMOL

Model	Total Train Time
Adapter-Transformer	44 mins 10 secs
LAMOL	48 mins 39 secs

Table 4: Training time taken for our proposed Adapter-Transformer method and baseline LAMOL

6 Conclusion

We propose Adapter-Transformers, a framework for LLL, that can contain information of new tasks without forgetting older tasks. With the current design, we can improve average task accuracy by 1.3% over LAMOL without increasing the network growth by more than 15% and avoiding replay and task order biasness. More tasks can be added as needed.

Limitations

As the number of tasks increases, the parameters can keep growing, although the growth is nominal. In our future work, we would like to consider the similarity in tasks to reduce the parameter growth. Also, we would like to study the performance of the proposed approach when number of tasks grows large.

References

- Luheng He, Mike Lewis, and Luke Zettlemoyer. 2015. [Question-answer driven semantic role labeling: Using natural language to annotate natural language](#). In *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing*, pages 643–653, Lisbon, Portugal. Association for Computational Linguistics.
- Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski,




Figure 3: Loss for the proposed model for the set of three NLP tasks

Bruna Morrone, Quentin de Laroussilhe, Andrea Gesmundo, Mona Attariyan, and Sylvain Gelly. 2019. [Parameter-Efficient Transfer Learning for NLP](#). Technical Report arXiv:1902.00751, arXiv. ArXiv:1902.00751 [cs, stat] type: article.

Yufan Huang, Yanzhe Zhang, Jiaao Chen, Xuezhi Wang, and Diyi Yang. 2021. [Continual Learning for Text Classification with Information Disentanglement Based Regularization](#). Technical Report arXiv:2104.05489, arXiv. ArXiv:2104.05489 [cs] type: article.

Kasidis Kanwatchara, Thanapapas Horsuwan, Piyawat Lertvittayakumjorn, Boonserm Kijsirikul, and Peerapon Vateekul. 2021. [Rational LAMOL: A Rationale-based Lifelong Learning Framework](#). In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 2942–2953, Online. Association for Computational Linguistics.

Ilya Loshchilov and Frank Hutter. 2019. [Decoupled Weight Decay Regularization](#). Technical Report arXiv:1711.05101, arXiv. ArXiv:1711.05101 [cs, math] type: article.

Bryan McCann, Nitish Shirish Keskar, Caiming Xiong, and Richard Socher. 2018. [The Natural Language Decathlon: Multitask Learning as Question Answering](#). Technical Report arXiv:1806.08730, arXiv. ArXiv:1806.08730 [cs, stat] type: article.

Nikola Mrkšić, Diarmuid Ó Séaghdha, Tsung-Hsien Wen, Blaise Thomson, and Steve Young. 2017. [Neural belief tracker: Data-driven dialogue state tracking](#). In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1777–1788, Vancouver, Canada. Association for Computational Linguistics.

- Jonas Pfeiffer, Andreas Rücklé, Clifton Poth, Aishwarya Kamath, Ivan Vulić, Sebastian Ruder, Kyunghyun Cho, and Iryna Gurevych. 2020. [Adapterhub: A framework for adapting transformers](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP 2020): Systems Demonstrations*, pages 46–54, Online. Association for Computational Linguistics.
- Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. 2018. [Language models are unsupervised multitask learners](#).
- Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher D. Manning, Andrew Ng, and Christopher Potts. 2013. [Recursive deep models for semantic compositionality over a sentiment treebank](#). In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*, pages 1631–1642, Seattle, Washington, USA. Association for Computational Linguistics.
- Fan-Keng Sun, Cheng-Hao Ho, and Hung-Yi Lee. 2019. [LAMOL: LAnguage MOdeling for Lifelong Language Learning](#). Technical Report arXiv:1909.03329, arXiv. ArXiv:1909.03329 [cs] type: article.