

VL-ADAPTER: Parameter-Efficient Transfer Learning for Vision-and-Language Tasks

Yi-Lin Sung Jaemin Cho Mohit Bansal
 UNC Chapel Hill

{yilsung, jmincho, mbansal}@cs.unc.edu

Abstract

Recently, fine-tuning language models pre-trained on large text corpora have provided huge improvements on vision-and-language (V&L) tasks as well as on pure language tasks. However, fine-tuning the entire parameter set of pre-trained models becomes impractical since the model size is growing rapidly. Hence, in this paper, we introduce adapter-based parameter-efficient transfer learning techniques to V&L models such as VL-BART and VL-T5. We evaluate our methods in a unified multi-task setup on both image-text and video-text benchmarks. For the image-text tasks, we use four diverse V&L datasets: VQAv2, GQA, NLVR², and MSCOCO image captioning. For video-text tasks, we use TVQA, How2QA, TVC, and YC2C. With careful training and thorough experiments, we benchmark three popular adapter-based methods (Adapter, Hyperformer, Compacter) against the standard full fine-tuning and the recently proposed prompt-tuning approach. We also enhance the efficiency and performance of adapters by sharing their weights to attain knowledge across tasks. Our results demonstrate that training the adapter with the weight-sharing technique (4.18% of total parameters for image-text tasks and 3.39% for video-text tasks) can match the performance of fine-tuning the entire model. Lastly, we present a comprehensive analysis including the combination of adapter and task-specific prompts and the impact of V&L pre-training on adapters.¹

1. Introduction

Following the success in the language domain [4, 8, 25, 30, 39, 40], large-scale pre-training of vision-and-language (V&L) models has become a standard framework to tackle V&L tasks [6, 7, 18, 32, 45, 48]. In such frameworks, V&L models, which are usually the combination of vision encoders and language models, are first pre-trained on large-

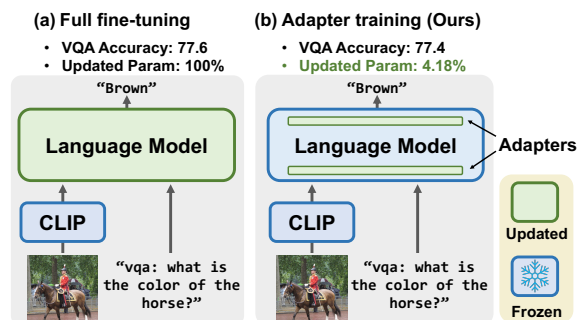


Figure 1. Comparison of (a) full fine-tuning and our (b) adapter training for V&L tasks. By updating only a small set of adapter parameters, we can achieve similar performance to full fine-tuning. We experiment with our adapter training on diverse image-text and video-text benchmarks, and here we show VQA as an example.

scale unlabeled data, then fine-tuned for downstream V&L tasks. This is the standard strategy to fuse the knowledge of vision-and-language to the language model. However, given that such models’ size grows very rapidly nowadays, either pre-training or fine-tuning of the V&L model can still contribute to an unignorable, large memory and storage cost. For instance, if we use GPT-3 [4] with 175B parameters as a backbone of V&L model, we would need 700 GB of memory to store its entire parameters.² To address this problem, recently, several parameter-efficient training methods have been proposed [12, 16, 17, 20, 28, 33, 47, 52]. Among them, adapter [16] and its variants [20, 33] are widely used in the NLP domain and applied to different architectures. Adapter is a small module added to intermediate layers of the model (which is illustrated in Figure 2), which allows to achieve as high performance as full fine-tuning (i.e., updating all parameters), by fine-tuning only a small set of parameters. Moreover, this also shows that it is possible to use a few parameters to learn the information fusion of vision and language without losing performance. Despite adapters having achieved success in text classification [16, 20, 33] and image-text alignment [2], to the best of

¹The code for our CVPR 2022 paper is available at: https://github.com/yilsung/VL_adapter.

² $(175 \times 10^9) \times 4(\text{bytes}) \times \frac{1}{10^{-9}}(\text{GB/bytes}) = 700(\text{GB})$

our knowledge, no work has utilized this efficient method for more challenging downstream V&L problems, such as visual/video question answering and image/video captioning. Besides, V&L models often come with expensive computations by combining the knowledge of two input modalities. Hence, we investigate the application of adapter-based parameter-efficient training techniques to V&L tasks.

We aim to efficiently tune language models on diverse downstream V&L tasks while achieving performance comparable to full fine-tuning. For this, we analyze these parameter-efficient training techniques in a unified multi-task learning setup, and we benchmark different adapter [16, 20, 33] and prompt-based methods [24]. For our V&L model, following Cho *et al.* [7], we adopt encoder-decoder language models (BART [25] and T5 [40]) that tackle multiple V&L tasks as text generation to avoid designing task-specific architectures. We use CLIP [37], a pretrained image-text alignment model, as our visual encoder for the ease of doing V&L pre-training. To inform the model about which task it is going to perform, we follow [7, 40] to add task-specific (text) prompts to the front of the input sentence (e.g., “vqa: [Q]” for VQA). We then insert Adapter [16] and its variants, Hyperformer [20] and Compacter [33], into the model to perform parameter-efficient training. Hyperformer and Compacter are recently proposed *state-of-the-art* approaches: Hyperformer improves the efficiency of adapters by generating their weights via a hyper-network, while Compacter reduces the parameters by utilizing Kronecker products and low-rank parameterization for the adapters’ weights. We also compare adapter-based approaches with prompt tuning [24], which adds trainable prompts to the input. We show the high-level concept of our work in Figure 1. Practically, adapter training involves parameter updates of adapter modules, layer normalization layers, and the visual projection layer (see Section 3.1 and Figure 2 for more details). Since we tackle multiple tasks simultaneously [20, 36], we also explore taking advantage of the sharing of information between tasks on adapters and prompts. Specifically, we make some of the trainable parameters to be shareable to learn cross-task information while reserving the rest of them for task-specific information. With this technique, the number of trainable parameters can be reduced even further.

We conduct our experiments and analysis on four diverse image-text tasks: VQAv2 [10], GQA [19], NLVR² [46], and MSCOCO captioning [5]. For completeness, we also verify the effectiveness of our framework on four video-text tasks: TVQA [22], How2QA [26], TVC [23], and YC2C [56]. Overall, the performance of the three adapter-based approaches closes the gap between which of full fine-tuning. In our experiments, Compacter does not stand out in terms of efficiency, since we remove the low-rank approximation for trading performance. Hyperformer is more

efficient than adapters, but we eventually show *our adapter training with the weight-sharing technique can achieve the same performance as full fine-tuning while only updating 4.18% of the entire parameters for image-text tasks (and 3.39% for video-text tasks)*. Next, we compare the fine-tuning and freezing of the CLIP parameters, where the latter achieves a better trade-off between performance and parameter efficiency. We also present a detailed analysis to understand the contribution of each component in adapters, as well as the different parameter-sharing mechanisms. We find that using a single set of adapter modules across all tasks achieves the best results and accuracy-efficiency trade-off, showing the possibility of pursuing efficiency with simplicity (Fig. 4). Since most of the experiments are based on the pre-trained weights accompanied with the models (e.g., CLIP pre-trained weights for ResNet and BART pre-trained weights), we also demonstrate that the results of training adapters on top of V&L pre-trained weights can match or even exceed which of the full fine-tuning counterpart. While we conduct most of our experiments with the V&L generation framework, we also extend the adapter training to CLIP-ViL [45], which is one of the SOTA discriminative V&L approaches. Lastly, we report the results of comprehensive hyperparameter search in Appendix, hoping that they will be useful for related research on parameter-efficient training.

Our contributions could be summarized as: (1) the first work benchmarking different types of parameter-efficient training techniques (Adapter, Hyperformer and Compacter) for various challenging downstream image-text and video-text tasks; (2) empirical demonstration of adapters reaching the performance of full fine-tuning while updating only 3.39-4.18% of the parameters; (3) comprehensive analysis on the design of freezing CLIP, impact of different architectural components, weight-sharing techniques, task-specific prompts, and vision-language pretraining.

2. Related Work

Our research is built upon previous works about language models, V&L models, and parameter-efficient training. In this section, we introduce previous literature and discuss their similarities and differences with this work.

Language Model Pre-training. The workflow of pre-training and fine-tuning has become a popular paradigm for solving many downstream tasks in the NLP field. Several papers accordingly propose new architectures and objectives to model language, such as ELMo [35], BERT [8], RoBERTa [31], GPT series [4, 38, 39], and encoder-decoder versions such as BART [25] and T5 [40]. Comprehensive studies [40] have shown the effectiveness of the encoder-decoder model compared to other architectures. Hence, we choose BART and T5 as our text generative model.

V&L Pre-training. To tackle V&L tasks, most existing approaches combine respective models specialized for either pure language or pure vision tasks as a V&L model. ViLBERT [32], UNITER [6] and LXMERT [48] use the faster R-CNN [43] object detector to extract the bottom-up [1] features from the image and provide them with text features to a cross-modality transformer to solve visual question answering and image-text alignment tasks. Cho *et al.* [7] adopt encoder-decoder language models [25, 40] to generalize V&L tasks to text generation with a task-agnostic architecture, which combines faster R-CNN object detector and T5 (or BART). PixelBERT [18] replaces R-CNN with plain ResNet [14], and demonstrates the advantage of including the vision model in the training. Inspired by the success of pre-training with web-scale unlabeled data in the NLP field, Radford *et al.* [37] propose CLIP, which has a rich cross-modal representation and can solve a wide range of tasks without additional supervision. CLIP-ViL [45] explores the advantage of using a CLIP visual encoder (ResNet or ViT [9]) for V&L tasks. Although language models are tuned for downstream tasks in previous works, some recent research [49, 51] attempts to freeze large language models (e.g., GPT-3) to achieve zero-shot learning for V&L tasks. This line of research focuses on how to map images to the inputs that the language model can use. Frozen [49] achieves this by jointly training an NF-ResNet-50 [3] and frozen GPT-3 with the Conceptual Captioning dataset [44]. Instead of aligning images to text features, Yang *et al.* [51] directly utilize a pre-trained image captioner to transform images to text, which is a useful resource for a language model, and they demonstrate the effectiveness of this framework on visual question answering tasks. Our V&L model is the combination of CLIP and BART (or T5). We conduct a thorough ablation study to test the performance of four training scenarios: all possible pairs of training or freezing CLIP and BART. The results in Section 5.2 show that training the BART model only has the best trade-off between performance and parameter efficiency, hence, even though the architecture is end-to-end trainable, we decide to freeze the CLIP to fulfill our goal of parameter-efficient training.

Parameter-Efficient Training. As the size of recent models increases rapidly, updating the models in parameter-efficient ways becomes crucial. Recently, three types of methods have been proposed: (1) only updating newly added parameters (added either to the input or model); [16, 20, 24, 28, 33]; (2) sparsely updating a small number of parameters of the model; and [12, 47, 52] (3) low-rank factorization for the weights to be updated [17]. [13, 34] combine such approaches to propose a unified parameter-efficient training framework. Among these approaches, adapters, which belong to the first category, are widely used in computer vision [41, 42] and natural language processing [16, 20, 33]. While adapters add additional parameters

into models, prompt-based approaches instead add trainable parameters into the inputs [11, 24, 28], and experiments have shown their value in language tasks. Some concurrent works also extend parameter-efficient techniques to CLIP models [2, 54, 55]. However, they mainly tackle image-text alignment problems, while in this paper we aim at more downstream V&L tasks, such as visual/video question answering, visual reasoning, and image/video captioning. The typical usage of adapters is training them independently on different tasks. This manner prevents these adapters from learning the shared information across tasks. In this paper, we find that making these plug-and-play adapters shareable improves the performance of a low resource dataset and reduces the overall trainable parameters.

3. Methods

Our main contribution is to exploit adapter-based methods to efficiently fine-tune generative models in the multi-task setting. We explore several *state-of-the-art* variations of adapters (Adapter, Hyperformer, Compacter) on various V&L tasks and show that vanilla adapters are the best among them. We further demonstrate that sharing adapters across tasks can boost performance to match full fine-tuning results and further improve the parameter efficiency.

3.1. Unified Framework for V&L Tasks

We illustrate our V&L model in Figure 2(a). We follow [7] to unify V&L tasks to a text generation problem. Our V&L model is a combination of CLIP and BART (T5), and therefore we name our base architecture CLIP-BART (CLIP-T5). To be more specific, assuming that we have a pair of an image (or video) \mathbf{x}^I and a sentence \mathbf{x}^S (e.g. the question in VQA) as the input for our model, our goal is to maximize the agreement between the model’s output and the text label of M tokens $\mathbf{y} = (y_1, y_2, \dots, y_M)$ (e.g., the ground truth answer in VQA). Regarding architectures, we use an encoder-decoder language model (parameterized by θ^L) as our main generative model. We connect a CLIP (parameterized by θ^V) and a visual projection layer (parameterized by $\theta^{V \rightarrow L}$; it projects the visual representation to the correct dimension for language model) to the model for extracting the visual representation from input images and feed the concatenation of visual representation and sentence representation to the encoder-decoder model. We train our model with cross-entropy (CE) loss:

$$\begin{aligned} l(\mathbf{x}^I, \mathbf{x}^S, \mathbf{y}; \theta^L, \theta^V, \theta^{V \rightarrow L}) \\ = - \sum_{i=1}^M y_i \log(f_{\theta}(\mathbf{x}^{V \rightarrow L}, \mathbf{x}^S)_i) \end{aligned} \quad (1)$$

where f_{θ} and $\mathbf{x}^{V \rightarrow L}$ refer to our model and the projected visual representation, respectively.

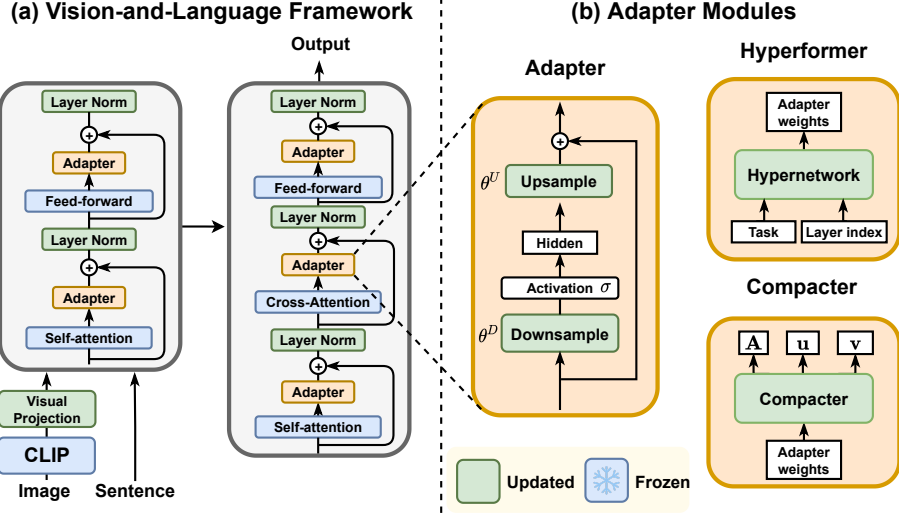


Figure 2. Illustration of (a) our unified framework for V&L tasks (Sec. 3.1) and (b) three-adapter based modules (Sec. 3.2). Green color refers to trainable parameters and blue color refers to frozen ones.

We next introduce our multi-task setup for the unified generative model. Benefiting from the unified format, we can construct a universal dataset \mathcal{D} from N given V&L datasets, $\mathcal{D}_1, \mathcal{D}_2, \dots, \mathcal{D}_N$. Thus, we optimize our parameters by minimizing the average loss over \mathcal{D} :

$$\mathcal{L}(\mathcal{D}; \theta^L, \theta^V, \theta^{V \rightarrow L}) = \frac{1}{|\mathcal{D}|} \sum_{(x^I, x^S, y) \in \mathcal{D}} l(x^I, x^S, y; \theta^L, \theta^V, \theta^{V \rightarrow L}) \quad (2)$$

Our trainable parameters are $\{\theta^L, \theta^V, \theta^{V \rightarrow L}\}$. As deep learning models have grown rapidly in recent times, updating and storing the whole parameters of either visual or language models can be inefficient. Therefore, this motivates us to introduce adapter-based approaches into V&L models for parameter-efficient tuning.

3.2. Adapters for V&L Models

Adapters. Figure 2(b) left. Adapters [16] are sub-networks with small parameters that are inserted after every attention and feed-forward layer in a model. With adapters, the models learn downstream tasks by updating only a small number of parameters. The adapters consist of a pair of downsampling and upsampling layers, and a residual connection. To be more specific, we denote the input of the adapter as $\mathbf{x} \in \mathbb{R}^{d_i}$, and the weight matrices for downsampling and upsampling layers to be $\theta^D \in \mathbb{R}^{d_i \times d}$ and $\theta^U \in \mathbb{R}^{d \times d_i}$, where d_i and d are the input and hidden dimensions, respectively. The mechanism of adapters is defined as:

$$h = f_{\theta^U}(\sigma(f_{\theta^D}(\mathbf{x}))) + \mathbf{x} \quad (3)$$

where $\sigma(\cdot)$ is an activation function, and we use GELU [15] in this paper. With adapters, the parameter complexity (i.e., the number of added parameters) is $\mathcal{O}(d_i d)$, and it

usually is $2 \sim 3\%$ of the whole model’s parameters. Note that all layer normalization layers are also updated to adapt to the data distribution of downstream data.

Hyperformers. Figure 2(b) top-right. The typical usage of adapters is to separately train one adapter for one task. In that fashion, adapter modules are independent across tasks, preventing the possibility of reducing the required parameters by sharing the weights for similar tasks. Hence, in order to make the adapter module even more efficient, we extend the Hyperformer [20] to a V&L architecture. More specifically, we maintain a hyper-network that is shared over tasks to generate the adapters’ weights conditioned on the task and the index of the layer. Suppose that we have N_T tasks at hand, the number of layers of the model is N_L , and we can use $\mathbf{t}_1, \mathbf{t}_2, \dots, \mathbf{t}_{N_T} \in \mathbb{R}^{d_e}$ and $\mathbf{l}_1, \mathbf{l}_2, \dots, \mathbf{l}_{N_L} \in \mathbb{R}^{d_e}$ to represent their d_e dimensional embeddings. Hyperformer is composed of a two-layer task projector network $\theta^T \in \mathbb{R}^{d_e \times 2 \times d_p}$ and the hyper-network $\theta^H \in \mathbb{R}^{d_p \times (2 \times d \times d_i)}$, which aims to generate the weights for the upsampling and downsampling layer in adapters based on the projected embedding. Without loss of generality, to generate the adapter’s weights in the i^{th} layer for j^{th} task, the generation process is:

$$[\theta^D, \theta^U] = f_{\theta^H}(f_{\theta^T}([\mathbf{t}_j, \mathbf{l}_i])) \quad (4)$$

where $[\cdot]$ refers to concatenation. Note that to save memory with Hyperformer, the number of trainable parameters of Hyperformer needs to be smaller than that of adapters, namely, $|\theta^H| + |\theta^T| + N_T |\mathbf{t}| + N_L |\mathbf{l}| < N_T N_L (|\theta^U| + |\theta^D|)$. In general, we have $N_T |\mathbf{t}|, N_L |\mathbf{l}|, |\theta^T| \ll |\theta^H|$, so we can further induce the appropriate range of d_p , which is $d_p < N_T N_L$.

Compacters. Figure 2(b) bottom-right. Although adapters




Figure 3. Illustration of different weight sharing methods for adapters. Parameters with same color are shared across tasks.

have attained great success on parameter-efficient training, they still have redundant parameters and usually underperform full fine-tuning. Compacter hence is introduced by [33] to solve the issues with the matrix decomposition and parameter sharing, and eventually, Compacter has been shown to have a better trade-off between performance and efficiency compared to adapters. In the following, we demonstrate the mechanism of Compacter with the weights of the downsampling layer. First, Compacter introduces *parameterized hypercomplex multiplication layers* (PHM layers) [53], whose parameters are the decomposition of $\theta^D \in \mathbb{R}^{d_i \times d}$ to the sum of k Kronecker products:

$$\theta^D = \sum_{i=1}^k A_i \otimes B_i \quad (5)$$

where $A_i \in \mathbb{R}^{k \times k}$, $B_i \in \mathbb{R}^{\frac{d_i}{k} \times \frac{d}{k}}$. The parameter complexity of the PHM layer is $\mathcal{O}(\frac{d_i d}{k})$, reducing the cost by at most $\frac{1}{k}$. To further improve the efficiency of PHM layers, Compacter shares the parameters of a smaller matrix A_i across all layers, and decomposes the bigger matrix B_i even more with low-rank parameterization. Specifically, the matrix B_i is approximated to be a low-rank matrix, which is the product of two low-rank matrices, $\mathbf{u}_i \in \mathbb{R}^{\frac{d_i}{k} \times r}$ and $\mathbf{v}_i \in \mathbb{R}^{r \times \frac{d}{k}}$, where r is the matrix’s rank. This results in *low-rank parameterized hypercomplex multiplication layer* (LPHM):

$$\theta^D = \sum_{i=1}^k A_i \otimes B_i = \sum_{i=1}^k A_i \otimes (\mathbf{u}_i \mathbf{v}_i) \quad (6)$$

Empirically, $r = 1$ is sufficient to achieve competitive performance, obtaining the complexity of the LPHM layer as $\mathcal{O}(\frac{d+d_i}{k})$. Nevertheless, we find sharing matrix and low-rank decomposition in LPHM layers hurt the performance of V&L tasks, so we remove them and use the PHM layers instead.

Shared-Weight Adapters. Figure 3. Inspired by Hyperformer, we explore the sharing of information between N_T

tasks in vanilla adapters with the weight-sharing technique. We denote $\Theta = \{\Theta^D, \Theta^U\}$, is the collection of all inserted adapter modules’ weights in the model, and Θ^D (Θ^U) is the subset that only includes downsampling (upsampling) layers in adapters. As we have mentioned earlier, adapters are trained independently, so we have unique $\{\Theta_i^D, \Theta_i^U\}$ for the i^{th} task. To enable the adapter to learn cross-task information, we make part of the weights of the adapter to be shareable. For instance, we can make Θ_i^D equal to Θ_j^D ($i \neq j$), and the rest of parameters (Θ_i^U) can still learn the task-specific information for the i^{th} task. Note that we also consider the extreme case of using a single adapter for all tasks ($\Theta_i = \Theta_j$). We illustrate different weight sharing methods for adapters in Figure 3.

Where to Add Adapters? Our goal is to apply adapters into V&L models to efficiently update their parameters. Since our architecture is composed of visual and language components, it is expected that adapter layers are injected into both of them. However, we observe that even fully fine-tuning the whole model does not bring much improvement compared to only updating the language one (see details in Section 5.2). Since freezing CLIP has the better trade-off between performance and parameter efficiency over training it, we eventually do not add adapter layers into CLIP.

Multi-task Adapter Variants. We consider several approaches to using adapters. Since we are in a multi-task setting, the first two straightforward methods are training adapters and Compacter per task, and we call **Multiple Adapters** and **Multiple Compacters** (illustrated in Figure 3(a)). To allow adapters to learn information across tasks, we use weight sharing techniques mentioned in Section 3.2 to form **Half-shared Adapters**, illustrated in Figure 3(b). Notably, we can form the two types of Half-shared Adapters by sharing either upsampling layers or downsampling layers. However, we use Half-shared Adapters to represent adapters with sharing upsampling layers, since they have almost no differences in terms of accuracy and efficiency. We also consider the extreme case of training multiple tasks with one set of adapter layers, and this brings **Single Adapter** and **Single Compacter** (see Figure 3(c) for details). Lastly, we have **Hyperformer** which essentially shares information from multiple tasks. Note that we always update $\theta^{V \rightarrow L}$ (1.14% of parameters) and all layer normalizations (0.04% of parameters) in the language model. We freeze the output layer, whose weights are tied with word embeddings, because it occupies about 30% of the language model’s parameters, and updating it doesn’t come with a performance boost.

4. Experimental Setup

Datasets. For image-text experiments, we evaluate our models on four V&L datasets: VQAv2 [10] and GQA [19]

Method	Updated Params (%)	VQA Karpathy test Acc. (%)	GQA test-dev Acc. (%)	NLVR ² test-P Acc. (%)	COCO Cap. Karpathy test CIDEr	Avg.
VL-BART [7]	100.00	67.8	57.3	72.3	109.4	76.7
CLIP-BART						
+ Full fine-tuning	100.00	67.6	56.7	73.0	112.9	77.6
+ Multiple Adapters	12.22	65.4	54.0	69.8	114.3	75.9
+ Half-shared Adapters	8.36	65.2	53.4	71.2	113.7	75.9
+ Single Adapter	4.18	65.9	54.5	74.2	114.9	77.4
+ Hyperformer	5.79	65.1	53.4	72.3	114.6	76.4
+ Multiple Compacters	7.05	64.6	53.4	69.1	116.0	75.8
+ Single Compacter	2.70	64.2	53.3	71.7	114.1	75.8
+ Multiple LoRA	17.72	65.5	53.0	62.8	115.4	74.2
+ Single LoRA	5.93	65.2	53.6	71.9	115.3	76.5
+ Multiple Prompts	4.53	43.8	38.1	51.1	104.6	59.4
+ Single Prompt	2.00	44.0	36.3	51.8	103.9	59.0

Table 1. The multi-task evaluation results on VQA, GQA, NLVR² and COCO Caption between full fine-tuning, adapter-based approaches, prompt-tuning, LoRA, and VL-BART. We bold the highest scores separately for approaches that are with or without parameter-efficient training techniques. We also report the results of the test-dev split on VQA in Appendix, and the trend is similar to using the Karpathy test set. Note that we don’t use V&L pre-training for every model.

for visual question answering, NLVR² [46] for visual reasoning, and MSCOCO [5] for image captioning. As for video-text experiments, we apply our method on four tasks from VALUE [27] benchmark: TVQA [22] and How2QA [26] for video question answering, TVC [23] and YC2C [56] for video captioning. The statistics of each dataset are shown in Appendix.

Architecture Details. We follow [7] to combine the vision encoder and an encoder-decoder language model to deal with many tasks via the unified text generation framework. For image experiments, we use CLIP-ResNet101 as our vision encoder [37]. Input images are resized to 224×224 for the memory efficiency. We extract the 7×7 grid features produced by the last convolutional layer, and then apply adaptive maximum-pooling over the features for down-sampling then to 6×6 for a fair comparison to [7]. For video experiments, we use the features extracted by CLIP (ViT-B/32) following [27], where they uniformly sample one frame per second and concatenate the CLIP outputs of the sampled frames to form the visual input. We limit the maximum length of visual input to 64 for efficiency. Following [27], we also include subtitles as additional information. In addition, we choose BART_{base} [25] as our main encoder-decoder language model, but we also extend the studies to T5_{base} [40]. We use CLIP-BART and CLIP-T5 for representing these two V&L architectures. We report all the training and evaluation details in Appendix.

5. Results and Analysis

5.1. Multi-Task Parameter-Efficient Fine-Tuning

Next, we move on to our main experiments on applying parameter-efficient training techniques for V&L models. We note that the techniques are only used for the

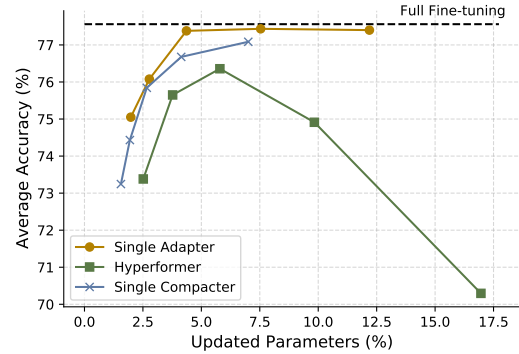


Figure 4. The comparison of three adapter-based approaches over the different percentages of updated parameters. We adjust the hidden dimension d to attain the model with different sizes.

language model, as we mentioned in Section 3.2. Besides the adapter-based methods and full fine-tuning, we also consider prompt-tuning [24] and LoRA [17], which are competitive parameter-efficient training approaches. In this case, we also have two variants for each method, that is, **Single Prompt-Tuning**, **Multiple Prompt-Tuning**, **Single LoRA** and **Multiple LoRA**, where the single one uses the same prompt/low-rank weights for multiple tasks while the multiple one has one prompt/low-rank weights for each task, respectively. The prompts are only added to the input of the encoder because we do not see improvement when the prompts are also added for the decoder. In the following paragraphs, we separately introduce the results in Table 1 and their takeaways. We also report the leaderboard results of multiple approaches on VQA and GQA in Appendix, where we select the representative approaches from each family of methods based on their performance in Table 1. The trend still holds similar to the results in Table 1.

Method	VQA	GQA	NLVR ²	COCO Cap.	Avg.
CLIP-T5					
+ Full fine-tuning	67.3	56.5	75.4	113.1	78.1
+ Single Adapter	67.6	56.2	73.9	111.8	77.4

Table 2. Results of Single Adapter & Full fine-tuning on CLIP-T5.

Single Adapter Performing the Best. We observe that the vanilla adapter is the most competitive architecture among the three types of adapter variants. We find Half-shared Adapters perform on par with Multiple Adapters with fewer parameters, and the Single Adapter’s performance is as competitive as which of the full fine-tuning (77.4 vs. 77.6). The performance boost of Half-shared Adapters and Single Adapter mainly comes from a smaller dataset, NLVR², and this demonstrates that information sharing benefits the low resource tasks.

We next turn to the results of Hyperformer and Compacter. The Hyperformer shares information across tasks in the hyper-network, thus resulting in that it is more parameter-efficient than the Multiple Adapters (12.22% vs. 5.79%). However, the Single Adapter still outperforms Hyperformer in terms of its parameter-efficiency (4.18% vs. 5.79%) and effectiveness (77.4 vs. 76.4). The optimization of hyper-network is harder and it might be one of the reasons to produce this outcome. Compared to Adapter, Compacter does not stand out in our experiments.

We hypothesize the reason causing the outcomes is: our BART model is pre-trained on pure language tasks, and we would like to adapt the model to perform on V&L tasks. Nevertheless, the assumption of Kronecker products might be too restrictive, so that Compacter fails to overcome the huge discrepancy between tasks. To have a complete comparison between three adapter-based approaches, in Figure 4, we show their performance over different percentages of updated parameters. We observe that Single Adapter, despite its simple architecture, achieves the best accuracy-efficiency trade-off.

Lastly, we transfer the best configuration of the single adapter to CLIP-T5 and show the results in Table 2. Note that we use a larger hidden dimension for adapters in this case, and the percentage of updated parameters is 7.98%. The results conclude that the Single Adapter still achieves a promising accuracy-efficiency trade-off in T5. We leave the results of adding other approaches to T5 in Appendix.

LoRA and Prompt-tuning vs. Adapters. Compared to the Single Adapter, LoRA uses slightly more parameters and the accuracy drops by approximately 1%. However, it is still a competitive approach since its performance is better than other methods. In general, prompt-tuning does not perform well in our experiments. The reason might be similar to Compacter’s: because the pre-trained tasks and downstream tasks are dissimilar, the model cannot adapt to the distribution of new datasets with few parameters. Also,

	fine-tuning BART	frozen BART
fine-tuning CLIP	65.6	39.4
frozen CLIP	64.7	39.1

Table 3. The VQA results of all possible pairs of full fine-tuning or freezing the CLIP and BART. Note that we only train the visual project layer ($\theta^{V \rightarrow L}$) in “frozen CLIP + frozen BART”.

prompt-tuning is sensitive to some training configurations such as model size, prompt initialization, and pre-training methods [24]. We leave the improvement of prompt-tuning performance for future work.

VL-Adapter in Video-Language Understanding Tasks. Table 4 shows the performance of the representative approaches used in Appendix on TVQA, How2QA, TVC, and YC2C. We transfer the hyperparameter setup used in Table 1, but we divide the learning by 3 for all approaches for better performance. Single Adapter again attains the best results among all parameter-efficient methods and is on par with full fine-tuning. From the image-text and video-text experiments, we find that the Adapter is more stable than other approaches.

5.2. Training or Freezing Visual Encoder

Because the use of CLIP enables the entire model to be end-to-end trainable with stability [45], we conduct experiments to test the performance of four different training scenarios: all possible pairs of training (full fine-tuning) or freezing CLIP and BART. Every combination is trained and evaluated on VQA. We apply the training tricks³ used in [18, 45] and report the results in Table 3. The results show that there is only 1% improvement with adding the CLIP into training. Given this result, the advantage of adding adapters inside CLIP over keeping it frozen is limited. Therefore, we have determined to freeze the CLIP consistently to save memory, and this also ideally fits our purpose of training models efficiently. We also report the result of our version of “Frozen” [49]⁴, where we freeze BART and only fine-tune CLIP. However, the accuracy of 39.4 (in Table 3) is far from that of updating BART, suggesting that fine-tuning the language model is still crucial.

5.3. Ablations and Analysis

In the following paragraphs, we sequentially present our ablation studies and analysis on (1) the contribution of different modules. (2) adapters with task-specific prompts. (3) effect of V&L pre-training on adapters. (4) adapters on the discriminative V&L model. We put results for (2) in Appendix.

³We have our best jointly training model with a 1×10^{-6} learning rate and the SGD optimizer on vision encoder while 1×10^{-4} and the AdamW optimizer for the rest of the model.

⁴However, unlike Frozen using both captioning and VQA data, we only use the latter in our experiments.

Method	Updated Params (%)	TVQA test-pub Acc. (%)	How2QA test-pub Acc. (%)	TVC test-pub CIDEr	YC2C test-pub CIDEr	Avg.
CLIP-BART						
+ Full fine-tuning	100.00	76.3	73.9	45.7	154.0	87.4
+ Single Adapter	3.39	76.6	73.9	46.3	152.9	87.4
+ Single LoRA	5.17	75.5	72.9	44.6	140.9	83.4
+ Single Prompt	1.18	32.1	30.2	29.5	102.0	48.4

Table 4. The multi-task evaluation results on TVQA, How2QA, TVC and YC2C between full fine-tuning, Single Adapter, LoRA and prompt-tuning.

Method	VQA	GQA	NLVR ²	COCO Cap.	Avg.
$\theta^{V \rightarrow L}$ only	32.2	25.6	52.1	78.5	47.1
$\theta^{V \rightarrow L}$ + Layer Norm	49.5	40.1	52.4	109.6	62.9
Single Adapter	65.9	54.5	74.2	114.9	77.4

Table 5. Ablation on adding components while adapter training.

Method	VQA	GQA	NLVR ²	COCO Cap.	Avg.
VL-BART [7]	69.1	59.0	73.3	111.5	78.2
CLIP-BART					
+ Full fine-tuning	69.2	57.5	75.0	112.1	78.5
+ Single Adapter	69.4	58.1	73.7	115.7	79.2

Table 6. The fine-tuning results of full fine-tuning and Single Adapter after pre-training on V&L tasks first.

Contribution of Modules ($\theta^{V \rightarrow L}$, Layer Norm). Recall that our adapter training not only includes the adapters’ modules but layer normalization layers and the visual projection layer $\theta^{V \rightarrow L}$. To have a better understanding of the contribution to the accuracy of adapters, we conduct ablation studies to gradually add trainable modules and compare their results in Table 5. We observe that only updating $\theta^{V \rightarrow L}$ produces insufficient results, suggesting the need to update the language model partially. We do find a considerable improvement with updating layer normalization layers, but the accuracy is still much behind that of the full fine-tuning, and this result also displays the effectiveness of adapters. We note that this finding deviates from the conclusion in [2], where updating layer normalization layer is comparable or even better than training adapters inside CLIP for the image classification task.

Effect of V&L Pre-training on Adapters. Since this work is about low-cost training, most of our experiments are based on the weights without V&L pre-training. This results in a performance gap between the approaches in this paper and *state-of-the-art* ones. We therefore also explore whether adapter training can take advantage of V&L pre-trained weights. We follow [45] to pre-train on COCO [29] and VG [21] images with multi-modal language modeling, visual question answering, and image-text matching. We exclude the referring tasks (grounded captioning and visual grounding) because they need bounding box information,

Method	Updated Params (%)	VQA test-dev	SNLI-VE dev	GQA test-dev
Full Fine-tuning	100.0	76.9	80.7	61.8
Single Adapter	4.3 - 6.2	76.5	80.2	60.7

Table 7. The results of adding adapters to CLIP-ViL on VQA, SNLI-VE, and GQA. Note that the number of parameters vary across tasks due to different output heads.

which cannot be obtained using CLIP. Refer to the training details in [45]. Table 6 shows the fine-tuning results with V&L pre-training. In this case, the Single Adapter even is more competitive than full fine-tuning with only a few parameters being trained, suggesting that the adapters work well with different kinds of pre-trained weights.

Adapters on the Discriminative V&L Pre-trained Model. While most of our experiments and ablations are conducted on generative V&L models, we also applied adapters to CLIP-ViL [45], a SOTA discriminative model, and report the results of VQA [10], SNLI-VE [50], and GQA [19] in Table 7. Adapters are also not added into CLIP due to the same reason in Section 5.2. Note that Single Adapter only updates from 4.3 to 6.2% of parameters with only a small gap in accuracy compared to full fine-tuning, which demonstrates the effectiveness of adapters for discriminative architectures.

6. Discussion and Conclusion

We conduct comprehensive studies on the evaluation of three adapter-based approaches on challenging image-text and video-text tasks. We employ a unified format and architecture to solve the tasks in a multi-tasking learning setup. With a thorough hyper-parameter search, we benchmark the performance of those methods and find the Single Adapter, which is a shared-weight vanilla adapter, is the best in terms of accuracy, efficiency, and simplicity. We have also shown that Single Adapter works well with a state-of-the-art discriminative model (CLIP-ViL). Lastly, we conduct ablation studies on understanding the contribution of different trainable modules, adapters with task-specific prompts, and the effect of V&L pre-training on adapters.

Acknowledgments

We thank the reviewers, Hyounghun Kim, Gedas Bertasius, and Hao Tan for their helpful discussions. This work was supported by ARO Award W911NF2110220, and ONR Grant N000141812871, and Google Focused Research Award. The views, opinions, and/or findings contained in this article are those of the authors and not of the funding agency.

References

- [1] Peter Anderson, Xiaodong He, Chris Buehler, Damien Teney, Mark Johnson, Stephen Gould, and Lei Zhang. Bottom-up and top-down attention for image captioning and visual question answering. *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 6077–6086, 2018. [3](#)
- [2] Anonymous. How to adapt your large-scale vision-and-language model. In *Submitted to The Tenth International Conference on Learning Representations*, 2022. under review. [1](#), [3](#), [8](#)
- [3] Andy Brock, Soham De, Samuel L Smith, and Karen Simonyan. High-performance large-scale image recognition without normalization. In *Proceedings of the 38th International Conference on Machine Learning*, pages 1059–1071. PMLR, 18–24 Jul 2021. [3](#)
- [4] Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language models are few-shot learners. *CoRR*, abs/2005.14165, 2020. [1](#), [2](#)
- [5] Xinlei Chen, Hao Fang, Tsung-Yi Lin, Ramakrishna Vedantam, Saurabh Gupta, Piotr Dollár, and C. Lawrence Zitnick. Microsoft COCO captions: Data collection and evaluation server. *CoRR*, abs/1504.00325, 2015. [2](#), [6](#)
- [6] Yen-Chun Chen, Linjie Li, Licheng Yu, Ahmed El Kholy, Faisal Ahmed, Zhe Gan, Yu Cheng, and Jingjing Liu. Uniter: Universal image-text representation learning. In *ECCV*, 2020. [1](#), [3](#)
- [7] Jaemin Cho, Jie Lei, Hao Tan, and Mohit Bansal. Unifying vision-and-language tasks via text generation. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139, pages 1931–1942. PMLR, 18–24 Jul 2021. [1](#), [2](#), [3](#), [6](#), [8](#)
- [8] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. In *NAACL*, 2019. [1](#), [2](#)
- [9] Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, and Neil Houlsby. An image is worth 16x16 words: Transformers for image recognition at scale. In *International Conference on Learning Representations*, 2021. [3](#)
- [10] Yash Goyal, Tejas Khot, Douglas Summers-Stay, Dhruv Batra, and Devi Parikh. Making the v in vqa matter: Elevating the role of image understanding in visual question answering. *2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 6325–6334, 2017. [2](#), [5](#), [8](#)
- [11] Yuxian Gu, Xu Han, Zhiyuan Liu, and Minlie Huang. Ppt: Pre-trained prompt tuning for few-shot learning. *ArXiv*, abs/2109.04332, 2021. [3](#)
- [12] Demi Guo, Alexander Rush, and Yoon Kim. Parameter-efficient transfer learning with diff pruning. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 4884–4896, Online, Aug. 2021. Association for Computational Linguistics. [1](#), [3](#)
- [13] Junxian He, Chunting Zhou, Xuezhe Ma, Taylor Berg-Kirkpatrick, and Graham Neubig. Towards a unified view of parameter-efficient transfer learning. *CoRR*, abs/2110.04366, 2021. [3](#)
- [14] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 770–778, 2016. [3](#)
- [15] Dan Hendrycks and Kevin Gimpel. Bridging nonlinearities and stochastic regularizers with gaussian error linear units. *CoRR*, abs/1606.08415, 2016. [4](#)
- [16] Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin De Laroussilhe, Andrea Gesmundo, Mona Attariyan, and Sylvain Gelly. Parameter-efficient transfer learning for nlp. In *International Conference on Machine Learning*, pages 2790–2799. PMLR, 2019. [1](#), [2](#), [3](#), [4](#)
- [17] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models. *CoRR*, abs/2106.09685, 2021. [1](#), [3](#), [6](#)
- [18] Zhicheng Huang, Zhaoyang Zeng, Bei Liu, Dongmei Fu, and Jianlong Fu. Pixel-bert: Aligning image pixels with text by deep multi-modal transformers. *CoRR*, abs/2004.00849, 2020. [1](#), [3](#), [7](#)
- [19] Drew A Hudson and Christopher D Manning. Gqa: A new dataset for real-world visual reasoning and compositional question answering. *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019. [2](#), [5](#), [8](#)
- [20] Rabeeh Karimi Mahabadi, Sebastian Ruder, Mostafa Dehghani, and James Henderson. Parameter-efficient multi-task fine-tuning for transformers via shared hypernetworks. In *Annual Meeting of the Association for Computational Linguistics*, 2021. [1](#), [2](#), [3](#), [4](#)
- [21] Ranjay Krishna, Yuke Zhu, Oliver Groth, Justin Johnson, Kenji Hata, Joshua Kravitz, Stephanie Chen, Yannis Kalantidis, Li-Jia Li, David A. Shamma, Michael S. Bernstein, and Li Fei-Fei. Visual genome: Connecting language and vision using crowdsourced dense image annotations. *International Journal of Computer Vision*, 123:32–73, 2016. [8](#)
- [22] Jie Lei, Licheng Yu, Mohit Bansal, and Tamara L. Berg. Tvqa: Localized, compositional video question answering. In *EMNLP*, 2018. [2](#), [6](#)
- [23] Jie Lei, Licheng Yu, Tamara L. Berg, and Mohit Bansal. Tvr: A large-scale dataset for video-subtitle moment retrieval. *ECCV*, 2020. [2](#), [6](#)
- [24] Brian Lester, Rami Al-Rfou, and Noah Constant. The power of scale for parameter-efficient prompt tuning. In *EMNLP*, 2021. [2](#), [3](#), [6](#), [7](#)

- [25] Mike Lewis, Yinhan Liu, Naman Goyal, Marjan Ghazvininejad, Abdelrahman Mohamed, Omer Levy, Veselin Stoyanov, and Luke Zettlemoyer. BART: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 7871–7880, Online, July 2020. Association for Computational Linguistics. 1, 2, 3, 6
- [26] Linjie Li, Yen-Chun Chen, Yu Cheng, Zhe Gan, Licheng Yu, and Jingjing Liu. Hero: Hierarchical encoder for video+language omni-representation pre-training. *EMNLP*, 2020. 2, 6
- [27] Linjie Li, Jie Lei, Zhe Gan, Licheng Yu, Yen-Chun Chen, Rohit Pillai, Yu Cheng, Luowei Zhou, Xin Eric Wang, William Yang Wang, et al. Value: A multi-task benchmark for video-and-language understanding evaluation. In *35th Conference on Neural Information Processing Systems (NeurIPS 2021) Track on Datasets and Benchmarks*, 2021. 6
- [28] Xiang Lisa Li and Percy Liang. Prefix-tuning: Optimizing continuous prompts for generation. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 4582–4597, Online, Aug. 2021. Association for Computational Linguistics. 1, 3
- [29] Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *European Conference on Computer Vision*, pages 740–755. Springer, 2014. 8
- [30] Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. Roberta: A robustly optimized BERT pretraining approach. *CoRR*, abs/1907.11692, 2019. 1
- [31] Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. Roberta: A robustly optimized BERT pretraining approach. *CoRR*, abs/1907.11692, 2019. 2
- [32] Jiasen Lu, Dhruv Batra, Devi Parikh, and Stefan Lee. Vilbert: Pretraining task-agnostic visiolinguistic representations for vision-and-language tasks. In *NeurIPS*, 2019. 1, 3
- [33] Rabeeh Karimi mahabadi, James Henderson, and Sebastian Ruder. Compacter: Efficient low-rank hypercomplex adapter layers. In *Thirty-Fifth Conference on Neural Information Processing Systems*, 2021. 1, 2, 3, 5
- [34] Yuning Mao, Lambert Mathias, Rui Hou, Amjad Almahairi, Hao Ma, Jiawei Han, Wen-tau Yih, and Madian Khabsa. Unipelt: A unified framework for parameter-efficient language model tuning. *CoRR*, abs/2110.07577, 2021. 3
- [35] Matthew E. Peters, Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, and Luke Zettlemoyer. Deep contextualized word representations. In *NAACL*, 2018. 2
- [36] Jonathan Pilault, Amine Elhattami, and Christopher Joseph Pal. Conditionally adaptive multi-task learning: Improving transfer learning in nlp using fewer parameters & less data. *ICLR*, 2021. 2
- [37] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In *ICML*, 2021. 2, 3, 6
- [38] Alec Radford and Ilya Sutskever. Improving language understanding by generative pre-training. In *arxiv*, 2018. 2
- [39] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, Ilya Sutskever, et al. Language models are unsupervised multitask learners. *OpenAI blog*, 1(8):9, 2019. 1, 2
- [40] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research*, 21(140):1–67, 2020. 1, 2, 3, 6
- [41] Sylvestre-Alvise Rebuffi, Hakan Bilen, and Andrea Vedaldi. Learning multiple visual domains with residual adapters. In *NIPS*, 2017. 3
- [42] Sylvestre-Alvise Rebuffi, Hakan Bilen, and Andrea Vedaldi. Efficient parametrization of multi-domain deep neural networks. *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 8119–8127, 2018. 3
- [43] Shaoqing Ren, Kaiming He, Ross B. Girshick, and Jian Sun. Faster r-cnn: Towards real-time object detection with region proposal networks. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 39:1137–1149, 2015. 3
- [44] Piyush Sharma, Nan Ding, Sebastian Goodman, and Radu Soricut. Conceptual captions: A cleaned, hypernymed, image alt-text dataset for automatic image captioning. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2556–2565, Melbourne, Australia, July 2018. Association for Computational Linguistics. 3
- [45] Sheng Shen, Liunian Harold Li, Hao Tan, Mohit Bansal, Anna Rohrbach, Kai-Wei Chang, Zhewei Yao, and Kurt Keutzer. How much can CLIP benefit vision-and-language tasks? *CoRR*, abs/2107.06383, 2021. 1, 2, 3, 7, 8
- [46] Alane Suhr, Stephanie Zhou, Ally Zhang, Iris Zhang, Hua-jun Bai, and Yoav Artzi. A corpus for reasoning about natural language grounded in photographs. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 6418–6428, Florence, Italy, July 2019. Association for Computational Linguistics. 2, 6
- [47] Yi-Lin Sung, Varun Nair, and Colin Raffel. Training neural networks with fixed sparse masks. In *Advances in Neural Information Processing Systems*, 2021. 1, 3
- [48] Hao Tan and Mohit Bansal. Lxmert: Learning cross-modality encoder representations from transformers. In *EMNLP*, 2019. 1, 3
- [49] Maria Tsimpoukelli, Jacob Menick, Serkan Cabi, S. M. Ali Eslami, Oriol Vinyals, and Felix Hill. Multimodal few-shot learning with frozen language models. *CoRR*, abs/2106.13884, 2021. 3, 7

- [50] Ning Xie, Farley Lai, Derek Doran, and Asim Kadav. Visual entailment: A novel task for fine-grained image understanding. *ArXiv*, abs/1901.06706, 2019. 8
- [51] Zhengyuan Yang, Zhe Gan, Jianfeng Wang, Xiaowei Hu, Yumao Lu, Zicheng Liu, and Lijuan Wang. An empirical study of gpt-3 for few-shot knowledge-based vqa. *ArXiv*, abs/2109.05014, 2021. 3
- [52] Elad Ben Zaken, Shauli Ravfogel, and Yoav Goldberg. Bitfit: Simple parameter-efficient fine-tuning for transformer-based masked language-models. *CoRR*, abs/2106.10199, 2021. 1, 3
- [53] Aston Zhang, Yi Tay, SHUAI Zhang, Alvin Chan, Anh Tuan Luu, Siu Hui, and Jie Fu. Beyond fully-connected layers with quaternions: Parameterization of hypercomplex multiplications with $\mathcal{O}(n)$ parameters. In *International Conference on Learning Representations*, 2021. 5
- [54] Renrui Zhang, Rongyao Fang, Peng Gao, Wei Zhang, Kunchang Li, Jifeng Dai, Yu Qiao, and Hongsheng Li. Tip-adapter: Training-free clip-adapter for better vision-language modeling. *ArXiv*, abs/2111.03930, 2021. 3
- [55] Kaiyang Zhou, Jingkang Yang, Chen Change Loy, and Ziwei Liu. Learning to prompt for vision-language models. *ArXiv*, abs/2109.01134, 2021. 3
- [56] Luowei Zhou, Chenliang Xu, and Jason J. Corso. Towards automatic learning of procedures from web instructional videos. In *AAAI*, 2018. 2, 6